

# Reinforcement Learning

## Lesson 4: Dynamic Programming

Edoardo Fazzari, 2022

# Overview

- Policy Evaluation (Prediction)
- Policy Improvement
- Policy Iteration
- Value Iteration
- Asynchronous Dynamic Programming
- Generalized Policy Iteration
- Efficiency of Dynamic Programming

# Dynamic Programming

## An introduction

- The term dynamic programming (DP) refers to a collection of algorithms that can be used to compute optimal policies given a perfect model of the environment as a Markov decision process (MDP)
- Classical DP algorithms are of limited utility in RL both because of their assumption and because of their great computational expense
- We assume that the environment is a finite MDP

# Policy Evaluation (Prediction)

# Policy Evaluation

## Iterative solution methods

- We already know the Bellman Equation
- If the environment's dynamic are completely known and the Bellman Equation is a system of  $|\mathcal{S}|$  simultaneous linear equation  $|\mathcal{S}|$  in unknowns
- For our purposes, *iterative solution methods* are the suitable
  - Consider a sequence of approximate value functions  $v_0, v_1, v_2, \dots$
  - The initial approximation  $v_0$  is chosen arbitrarily
  - Each successive approximation is obtained by using the Bellman equation as an update rule:

$$\begin{aligned} v_{k+1}(s) &\doteq \mathbb{E}_{\pi}[R_{t+1} + \gamma v_k(S_{t+1}) \mid S_t = s] \\ &= \sum_a \pi(a|s) \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_k(s')] \end{aligned} \tag{4.5}$$

# Iterative Policy Evaluation considerations

- $v_k = v_\pi$  is a fixed point for this update rule because the Bellman equation for  $v_\pi$  assures us of equality
- The sequence  $\{v_k\}$  can be shown in general to converge to  $v_\pi$  as  $k \rightarrow \infty$  under the same conditions that guarantee the existence of  $v_\pi$
- To produce each successive approximation,  $v_{k+1}$  from  $v_k$ , iterative policy evaluation applies the same operation to each state  $s$ : It replaces
  - the old value of  $s$  with a new value obtained from the old values of the successor state of  $s$
  - And, the expected immediate reward, along all the one-step transitions possible under the policy being evaluated (*Expected Update*)

# Iterative Policy Evaluation

## Pseudocode (estimating $V \approx v_\pi$ )

Input  $\pi$ , the policy to be evaluated

Algorithm parameter: a small threshold  $\theta > 0$  determining accuracy of estimation

Initialize  $V(s)$  arbitrarily, for  $s \in \mathcal{S}$ , and  $V(\text{terminal})$  to 0

Loop:

$\Delta \leftarrow 0$

    Loop for each  $s \in \mathcal{S}$ :

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until  $\Delta < \theta$

# Policy Improvement



# Policy Improvement Theorem

## Part 1

- Let's consider again the Bellman equation

$$\begin{aligned} q_{\pi}(s, a) &\doteq \mathbb{E}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_t = s, A_t = a] \\ &= \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_{\pi}(s')]. \end{aligned} \quad (4.6)$$

- The key criterion is whether this is greater than or less than  $v_{\pi}(s)$ . If it is greater:
  - that is, if it is better to select  $a$  once in  $s$  and thereafter follow  $\pi$  than it would be to follow  $\pi$  all the time
  - then one would expect it to be better still to select  $a$  every time  $s$  is encountered, and that the new policy would in fact be a better one overall

# Policy Improvement Theorem

## Part 2

- If that is true, we have a general result called *policy improvement theorem*
  - Let  $\pi$  and  $\pi'$  be any pair of deterministic policies such that, for all  $s \in \mathcal{S}$

$$q_{\pi}(s, \pi'(s)) \geq v_{\pi}(s) \quad (4.7)$$

- Then the policy  $\pi'$  must be as good as, or better than,  $\pi$ .
  - It must obtain greater or equal expected return from all states  $s \in \mathcal{S}$

$$v_{\pi'}(s) \geq v_{\pi}(s) \quad (4.8)$$

- **NOTE THAT:** If there is a strict inequality of (4.7) at any state, then there must be strict inequality of (4.8) at that state

# Greedy Policy

- So far we have seen how, given a policy and its value function, we can easily evaluate a change in the policy at a single state
- It is a natural extension to consider changes at *all* states, selecting at each state the action that appears best according to  $q_\pi(s, a)$  (*the greedy policy  $\pi'$* )

$$\begin{aligned}\pi'(s) &\doteq \operatorname{argmax}_a q_\pi(s, a) \\ &= \operatorname{argmax}_a \mathbb{E}[R_{t+1} + \gamma v_\pi(S_{t+1}) \mid S_t = s, A_t = a] \\ &= \operatorname{argmax}_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_\pi(s')]\end{aligned}\tag{4.9}$$

- By construction, it meets the conditions of the policy improvement theorem
  - so we know that it is as good as, or better than, the original policy

# Policy Improvement

## Definition

- The process of making a new policy that improves on an original policy, by making it greedy with respect to the value function of the original policy, is called *policy improvement*

# Optimal Policy

- Suppose the new greedy policy,  $\pi'$ , is as good as, but not better than, the old policy  $\pi$ . Then  $v_\pi = v_{\pi'}$
- From (4.9) it follows that for all  $s \in \mathcal{S}$ :

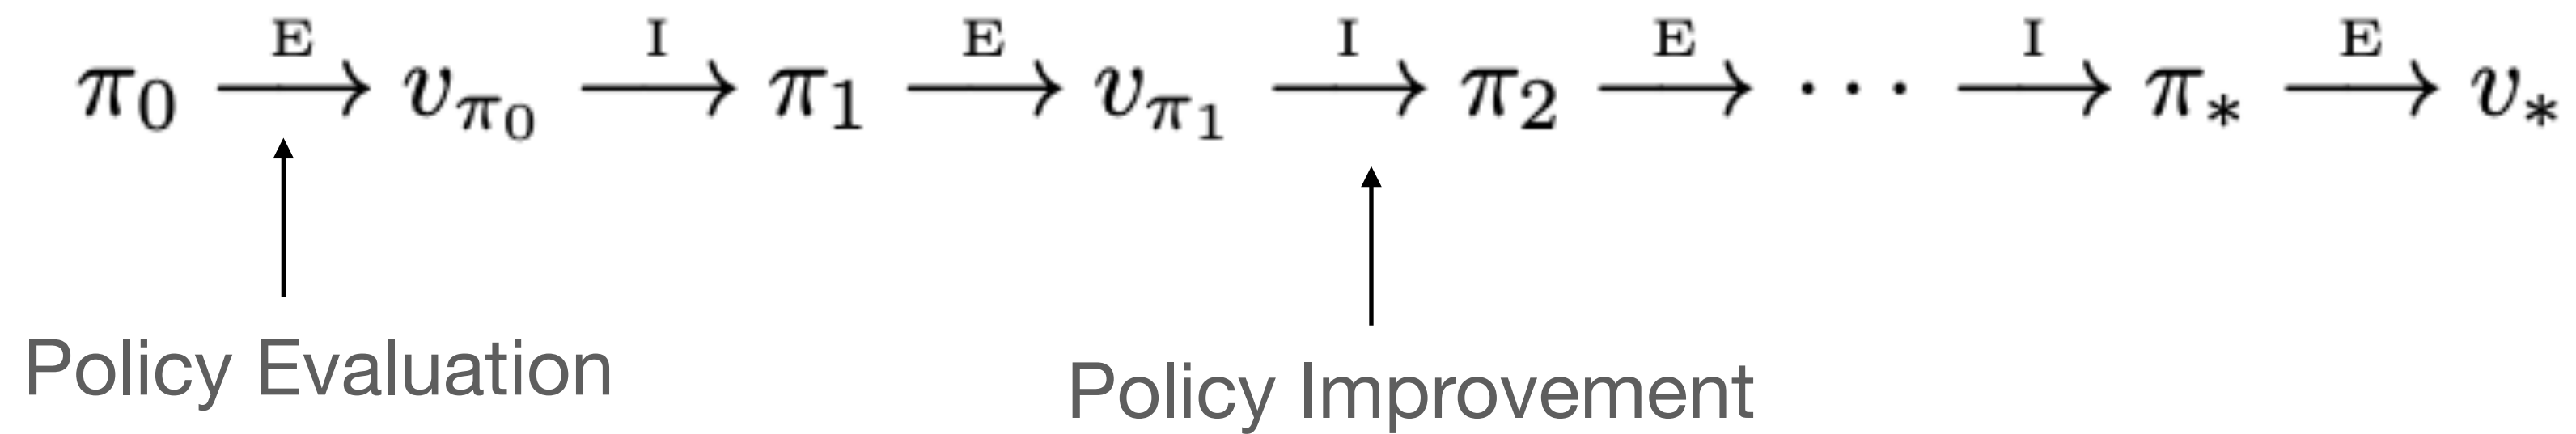
$$\begin{aligned} v_{\pi'}(s) &= \max_a \mathbb{E}[R_{t+1} + \gamma v_{\pi'}(S_{t+1}) \mid S_t = s, A_t = a] \\ &= \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_{\pi'}(s')]. \end{aligned}$$

- Which is the same of the Bellman optimality equation. Therefore:
  - $v'_\pi$  must be  $v_*$
  - $\pi$  and  $\pi'$  must be optimal policies

# Policy Iteration

# Policy Iteration

- Based on the previous results, a sequence of monotonically improving policies and value functions can be obtained:



- Because a finite MDP has only a finite number of deterministic policies, this process must converge to an optimal policy and the optimal value function in a finite number of iterations
- This way of finding an optimal policy is called *policy iteration*



# Policy Iteration

Using iterative policy evaluation for estimating  $\pi \approx \pi_*$

1. Initialization

$V(s) \in \mathbb{R}$  and  $\pi(s) \in \mathcal{A}(s)$  arbitrarily for all  $s \in \mathcal{S}$ ;  $V(\text{terminal}) \doteq 0$

2. Policy Evaluation

Loop:

$\Delta \leftarrow 0$

Loop for each  $s \in \mathcal{S}$ :

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_{s',r} p(s', r | s, \pi(s)) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until  $\Delta < \theta$  (a small positive number determining the accuracy of estimation)

3. Policy Improvement

*policy-stable*  $\leftarrow$  *true*

For each  $s \in \mathcal{S}$ :

*old-action*  $\leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$

If *old-action*  $\neq \pi(s)$ , then *policy-stable*  $\leftarrow$  *false*

If *policy-stable*, then stop and return  $V \approx v_*$  and  $\pi \approx \pi_*$ ; else go to 2



# Drawback

- Each iterations involves policy evaluation, which may itself be a protracted iterative computation requiring multiple sweeps through the state set
- If policy evaluation is done iteratively
  - Then convergence exactly to  $v_\pi$  occurs only in the limit
- *Must we wait for exact convergence, or can we stop short of that?*

# Value Iteration

# Value Iteration

## Part 1

- The policy evaluation step can be truncated in several ways without losing the convergence guarantees of policy iteration
- One important special case is when policy evaluation is stopped after just one sweep (one update of each state)
  - This algorithm is called *value iteration*

# Value Iteration

## Part 2

- This algorithm can be written as a simple update operation that combines the policy improvement and truncated policy evaluation steps:

$$\begin{aligned} v_{k+1}(s) &\doteq \max_a \mathbb{E}[R_{t+1} + \gamma v_k(S_{t+1}) \mid S_t = s, A_t = a] \\ &= \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_k(s')] \end{aligned} \tag{4.10}$$

- For all  $s \in \mathcal{S}$
- For arbitrary  $v_0$ , the sequence  $\{v_k\}$  can be shown to converge to  $v_*$ , under the same conditions that guarantee the existence of  $v_*$

# Value Iteration

For estimating  $\pi \approx \pi_*$

Algorithm parameter: a small threshold  $\theta > 0$  determining accuracy of estimation

Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$

Loop:

```
|  $\Delta \leftarrow 0$   
| Loop for each  $s \in \mathcal{S}$ :  
|    $v \leftarrow V(s)$   
|    $V(s) \leftarrow \max_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$   
|    $\Delta \leftarrow \max(\Delta, |v - V(s)|)$   
until  $\Delta < \theta$ 
```

Output a deterministic policy,  $\pi \approx \pi_*$ , such that

$$\pi(s) = \arg\max_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$$

# Value Iteration

## Final considerations

- Value iteration effectively combines, in each of its sweeps, one sweep of policy evaluation and one sweep of policy improvement
- Faster convergence is often achieved by interposing multiple policy evaluation sweeps between each policy improvement sweep

# Asynchronous Dynamic Programming

# Drawback of DP methods

- They involve operations over the entire state set of the MDP
  - They require sweeps of the state set
  - If the state set is very large, then even a single sweep can be prohibitively expensive



# Asynchronous DP

- Asynchronous DP algorithms are in-place iterative DP algorithms that are not organized in terms of systematic sweeps of the state set
- These algorithms update the values of states in any order whatsoever, using values of there states happen to be available
- The value of some states may be updated several times before the values of others are updated once
- To converge correctly, it must continue to update the values of all the states
- Asynchronous DP algorithms allow great flexibility in selecting states to update

# Avoiding sweeps

- Avoiding sweeps does not necessarily mean that we can get away with less computation
- It just means that an algorithm does not need to get locked into any hopeless long sweep before it can make progress improving a policy
- We can try to take advantage of this flexibility by selecting the states to which we apply updates so as to improve the algorithm's rate of progress
- We can try to order the updates to let value information propagate from state to state in an efficient way
  - Some states may not need their values updated as often as others
  - We might even try to skip updating states entirely if they are not relevant to optimal behavior


# Intermix computation

- Asynchronous algorithms make it easier to intermix computation with real-time interaction
- To solve a given MDP, we can run an iterative DP algorithm *at the same time that an agent is actually experiencing the MDP*
  - The agent's experience can be used to determine the state to which the DP algorithm applies its updates
  - At the same time, the latest value and policy information from the DP algorithm can guide the agent's decision making

# Generalized Policy Iteration

# Generalized Policy Iteration

- We use the term *generalized policy iteration (GPI)* to refer to the general idea of letting policy-evolution and policy-improvement processes interact, independent of the granularity and other details of the two process
- Almost all RL methods are well described as GPI:
  - All have identifiable policies and value functions,
  - With the policy always being improved with respect to the value function
  - and the value function always being driven toward the value function for the policy



# Competing and Cooperating

- The evaluation and improvement processes in GPI can be viewed as both competing and cooperating
- They compete in the sense that they pull in opposite directions
  - Making the policy greedy with respect to the value function typically makes the value function incorrect for the changed policy
  - Making the value function consistent with the policy typically causes that the policy no longer to be greedy
- In the long run, however, these two processes interact to find a single joint solution: the *optimal value function* and an *optimal policy*

# Efficiency of Dynamic Programming



# Some considerations

- DP may not be practical for very large problems, but compared with other methods for solving MDPs, DP methods are actually quite efficient
- If we ignore a few technical details, in the worst case, the time that DP methods take to find an optimal policy is *polynomial* in the number of states and actions
  - If they are started with good initial value functions or policies, they usually converge much faster than their theoretical worst-case run times
  - DP is comparatively better suited to handling large state spaces than competing methods such as *direct search* and *linear programming*
- In practice, DP methods can be used with today's computers to solve MDPs with millions of states
- Both policy iteration and value iteration are widely used, and it is not clear which, if either, is better in general



# Problems with large spaces

- On problems with large state spaces, *asynchronous DP methods* are often preferred
  - To complete even one sweep of a synchronous method requires computation and memory for every state
  - For some problems, even this much memory and computation is impractical
    - Yet the problem is still potentially solvable because relatively few states occur along optimal solution trajectories

# Bibliography:

Reinforcement Learning An Introduction (Second Edition), R. S. Sutton & A. G. Barto